



Bachelorarbeit

A Comparison of Synthetic-to-Real Domain Adaptation Techniques

Eberhard Karls Universität Tübingen
Mathematisch-Naturwissenschaftliche Fakultät
Wilhelm-Schickard-Institut für Informatik
Lernbasierte Computer Vision
Peter Trost, peter.trost@student.uni-tuebingen.de, 2019

Bearbeitungszeitraum: 24.05.2019-23.09.2019

Betreuer/Gutachter: Prof. Dr. Andreas Geiger, Universität Tübingen

Selbstständigkeitserklärung

Hiermit versichere ich, dass ich die vorliegende Bachelorarbeit selbständig und nur mit den angegebenen Hilfsmitteln angefertigt habe und dass alle Stellen, die dem Wortlaut oder dem Sinne nach anderen Werken entnommen sind, durch Angaben von Quellen als Entlehnung kenntlich gemacht worden sind. Diese Bachelorarbeit wurde in gleicher oder ähnlicher Form in keinem anderen Studiengang als Prüfungsleistung vorgelegt.

Peter Trost (Matrikelnummer 4039682), August 1, 2019

Abstract

Template

Acknowledgments

I want to thank everyone that helped me create this work. Especially Anastasia Bitter for her emotional and active support, Stephan Richter for explaining some aspects of the GTA dataset to me and Taesung Park for sending me a pretrained CycleGAN model. Further I want to thank my mentors Despoina Paschalidou and Yiyi Liao for their guidance and support and Professor Andreas Geiger for enabling me to work on this topic and the help in defining it precisely.

Contents

1. Introduction	11
2. Foundations	13
2.1. Domain Adaptation	13
2.2. Neural Networks	14
2.2.1. Convolutional Neural Networks	14
2.2.2. Generative Adversarial Networks	14
3. Related Work	17
3.1. section	17
4. Domain Adaptation Techniques	19
5. Experiments	21
5.1. Datasets	21
5.1.1. Synthetic dataset:	
Playing for Data: Ground Truth from Computer Games	21
5.1.2. Real dataset:	
The Cityscapes Dataset for Semantic Urban Scene Understanding	21
5.2. comparison Benchmark	22
5.2.1. Intersection over Union (IoU)	22
5.2.2. Methodology	22
6. Conclusion	23
A. Blub	25

1. Introduction

With autonomous driving being highly popular these days and the need for large amounts of labeled data to train the deep neural networks running the autonomous cars, methods have been developed to generate that data. Real datasets like Cityscapes [COR⁺16] or the Berkley Deep Drive [YXC⁺18] are expensive to assemble and even more expensive to label. Especially pixel-level annotations require a lot of working hours and can be inaccurate. In contrast synthetic approaches like the GTAV [RVRK16] or SYNTHIA datasets [RSM⁺16] enable to create vast amounts of image data automatically and make the annotation process much faster and more accurate and therefore cheaper than their real counterparts. The drawback of using synthetic data to train models that are ran in autonomous cars is that these models don't perform well in real environments. This can result in wrong predictions, which must be avoided in order to make autonomous driving possible. The drop in performance is due to the domain shift, i.e. changes in appearance, texture and others of objects in the synthetic images compared to objects in the real world seen through the car's cameras. An approach to create models to perform better in this scenario is domain adaptation. The idea behind that is to adapt a specific domain to another domain (e.g. synthetic to real, image taken at night to image taken at daytime, images of an object taken from one perspective to another perspective, ...). For autonomous driving this means making the synthetic image look more like an image of a real scene taken with a real camera. There are many techniques for adapting images from a synthetic to a realistic domain. Each having different approaches and using different methods. This work compares three current approaches on synthetic-to-real domain adaptation, namely CycleGAN [ZPIE17], CyCADA [HTP⁺17] and SG-GAN/Grad-GAN [LLJX18], all three of which are based on Generative Adversarial Networks [GPAM⁺14]. And tries to show their strengths and weaknesses when using them to translate images from the GTAV dataset [RVRK16] to images looking like the ones in Cityscapes [COR⁺16]. In order to evaluate the resulting images, a pre-trained deeplabv3 [CPSA17] model (code: [DLR]) is used to predict semantic segmentation maps and use the cityscapes evaluation code [CSR] to compute Intersection over Union (IoU) for these.

2. Foundations

This chapter will describe and explain the foundations necessary for this work. The term **Domain Adaptation** will be defined and described. Furthermore all relevant **Neural Network** architectures will be shown and explained. Specifically **Convolutional Neural Networks** and **Generative Adversarial Networks**.

2.1. Domain Adaptation

As described in [Csu17], Domain Adaptation is the task of transferring a machine learning model that is working well on a source data distribution to a related target data distribution. In this work we will focus on the adaptation from synthetic to real images. When talking about a synthetic image we are implying it was rendered from a virtual scene through a rendering engine like for example blender’s “Cycles” [Cycc] or graphics engine like Unity [Uni]. We define real images as taken from a real-world scene through some kind of camera. See Figure 2.1 for examples. In general there are many more domains, for example an image of a painting in the style of a particular artist, images of an object from different viewpoints, and many more.



Figure 2.1.: Example images for the the two domains relevant for this work. A real image from the cityscapes dataset [COR⁺16] (left) and a synthetic image from the GTA5 dataset [RVRK16] (right)

2.2. Neural Networks

2.2.1. Convolutional Neural Networks

see [Wu17] (Introduction to CNNs) Convolutional Neural Networks (CNNs) are Deep Neural Networks mostly used with images as input, consisting of convolutional layers, that extract features from input data (e.g. edges, curves, circles), pooling layers, commonly using max-pooling (mapping a subregion of the input to its maximum value) or average pooling (mapping the subregion to the average of the values) and a fully connected network often used for classification. **TODO: add more description and example images**

Convolutional Layer. In this layer a so called kernel iterates over the input. The kernel is a matrix of fixed size depending on the method used. A 3x3 kernel is commonly used. These kernels are used to extract features like diagonal lines, horizontal lines, curves and more. For example to extract a diagonal line a kernel could look like this:

$$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

Pooling Layer This layer reduces the size of the image by mapping an area to its maximum or average value in order to focus on the important features and save computing time.

Fully Connected Layer The last type of layers is a standard feed forward network that has one output node for each class the net has to be able to predict. It is used to classify the input of the net given the features extracted by the preceding layers.

TODO: example architectures VGG16, AlexNet,..?

2.2.2. Generative Adversarial Networks

Generative Adversarial Networks (GANs) implement a two-player-game: A Discriminator learns from a given data distribution what is “real”. The Generator generates data. The goal of the generator is to fool the discriminator into believing the generated data is “real” i.e. tries to create samples coming from the same distribution as the “real” data. The discriminator will label anything as “fake” that doesn’t resemble the learned “real” data distribution (2 classes classification, real or fake). This way GANs can learn to generate realistically looking images of faces, translate images of art from one style to another and improve semantic segmentation. The generative model generally uses *maximum likelihood estimation*. In [Goo17] it is described in the following way:

The basic idea of maximum likelihood is to define a model that provides an estimate of probability distribution, parameterized by parameters θ . We then refer to the **likelihood** as the probability that the model assigns to the training data: $\prod_{i=1}^m p_{\text{model}}(x^{(i)}; \theta)$

The parameter θ that maximizes the likelihood of the data is better found in log space [Goo17]

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^m p_{\text{model}}(x^{(i)}; \theta) \quad (2.1)$$

$$= \arg \max_{\theta} \log \prod_{i=1}^m p_{\text{model}}(x^{(i)}; \theta) \quad (2.2)$$

$$= \arg \max_{\theta} \sum_{i=1}^m \log p_{\text{model}}(x^{(i)}; \theta) \quad (2.3)$$

as the maximum of the function is at the same θ value and we now have a sum which as well eliminates the possibility of having underflow by multiplying multiple very small probabilities together. Formally GANs are a structured probabilistic model (more info: Chapter 16 of Goodfellow et al. 2016) containing latent variables z and observed variables x . The generator is defined by a function G that takes z as input and uses $\theta^{(G)}$ as parameters, the discriminator by a function D that takes x as input and uses $\theta^{(D)}$ as parameters. Both players have cost functions that are defined in terms of both players' parameters. The discriminator wishes to minimize $J^{(D)}(\theta^{(D)}, \theta^{(G)})$ and must do so by controlling only $\theta^{(D)}$. This is analogous for the generator: he tries to minimize $J^{(G)}(\theta^{(D)}, \theta^{(G)})$ while controlling only $\theta^{(G)}$. In contrast to an optimization problem that has a solution that is the (local) minimum (a point in parameter space where all neighboring points have greater or equal cost), the GAN objective is a game. The solution to a game is a Nash equilibrium [Nas50], meaning that each player chooses the best possible option or strategy in respect to what the other player(s) choose. Here the Nash equilibrium is a tuple $(\theta^{(D)}, \theta^{(G)})$ that is a local minimum of $J^{(D)}$ with respect to $\theta^{(D)}$ and a local minimum $J^{(G)}$ with respect to $\theta^{(G)}$. **The generator** is simply a differentiable function G . When z is sampled from some simple prior distribution, $G(z)$ yields a sample of x drawn from p_{model} . Typically a deep neural network is used to represent G .

The game plays out in two scenarios. The first is where the discriminator D is given random training examples x from the training set. The goal of the discriminator here is for $D(x)$ to be near 1. In the second scenario, inputs z to the generator are randomly sampled from the model's prior over the latent variables. The discriminator then receives input $G(z)$, a fake sample by the generator. The discriminator strives to make $D(G(z))$ approach 0 while the generator tries to make that approach 1. The game's Nash equilibrium corresponds to $G(z)$ being drawn from the same distribution as the training data. Under the assumption that the inputs to the discriminator are half real and half fake, this corresponds to $D(x) = \frac{1}{2}$ for all x .




Figure 2.2.: target data distribution of a toy dataset (right) and the data distribution of generated samples. Mind that the generated sample distribution hops from one mode to the next as the discriminator learns to recognize each mode as fake. Image from [MPPS16]

Training On each step, two minibatches are sampled: a minibatch of \mathbf{x} values from the dataset and one of \mathbf{z} values drawn from the model's prior over latent variables. Then two gradient steps are made simultaneously (simultaneous SGD): one updating $\theta^{(D)}$ to reduce $J^{(D)}$ and one updating $\theta^{(G)}$ to reduce $J^{(G)}$.

Cost functions There are many different cost functions that can be used for GANs. See chapter 4 for the ones that the techniques in this work use.

Advantages and Disadvantages

non-convergence GANs implement a two player game and in contrast to optimization problems that find a (local) minimum for example by stochastic gradient decent (SGD) and therefore usually make reliable downhill progress, this game has the following disadvantage: Updating one player and making the best current move downhill on the loss curve for that player might mean going uphill for the counterfeit player. Because of this GANs often oscillate in practice.

mode collapse happens when the generator learns to map different inputs \mathbf{z} to the same output. Partial mode collapse refers to scenarios in which the generator makes multiple images containing the same color or texture themes, or multiple images containing different views of the same dog. See Figure 2.2 for an example.

3. Related Work

There are many approaches of adapting Images from one domain to a different one. This chapter will show some of those techniques and compare them to the ones focused on in this work.

3.1. section

4. Domain Adaptation Techniques

This chapter will give an overview over the three **Domain Adaptation Techniques** that will be compared in this work, namely **CycleGAN** [ZPIE17], **CyCADA** [HTP⁺17] and **SG-GAN** [LLJX18]. It will explain how these techniques work, what kind of Network Architecture they use and how training these Nets was approached.

All of the methods compared in this work are based on Generative Adversarial Networks which were initially proposed in [GPAM⁺14] and are explained in Section 2. **CycleGAN** [ZPIE17] adds a cycle-consistency loss which consists of an additional generator/discriminator pair with the idea that translating a previously translated image back to the source domain should yield the original image again. Formally: let $G : X \rightarrow Y$ be the mapping function of the Generator translating an image from the source X to the target domain Y . Now add mapping function $F : Y \rightarrow X$ for the second generator's translation from target back to source domain. The goal of cycle-consistency is that for any image $x \in X$ from the source domain, $F(G(x)) \approx x$ holds. This analogously also has to hold for any image $y \in Y$ of the target domain with $G(F(y)) \approx y$. As described in chapter 2, GANs implement a two player game. The loss function of this game for CycleGAN is described as

$$\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [1 - \log D_Y(G(x))] \quad (4.1)$$

with D_Y being the discriminator that classifies samples $y \in Y$ as real or fake. This loss function is analogous for generator F and discriminator D_X . The cycle-consistency loss is defined as follow:

$$\mathcal{L}_{\text{cyc}}(G, F) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1] \quad (4.2)$$

with $\|\cdot\|_1$ being the L1-norm, which is the sum of absolute deviations between x and $F(G(x))$ and y and $G(F(y))$ with $x \in X$ and $y \in Y$, respectively. The full objective for the CycleGAN method is therefore:

$$\mathcal{L}(G, F, D_X, D_Y) = \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) + \lambda \mathcal{L}_{\text{cyc}}(G, F) \quad (4.3)$$

The discriminators' goal is to label the samples generated by the generators as fake while the generator tries to generate samples that the discriminator will label as real. This results in the goal of discriminators to maximize equation 4.2 and generators to minimize it. This leads to the goal for CycleGAN:

$$G^*, F^* = \arg \min_{G, F} \max_{D_X, D_Y} \mathcal{L}(G, F, D_X, D_Y) \quad (4.4)$$

CyCADA [HTP⁺17] adds a semantic loss to the CycleGAN approach to enforce semantic consistency (no more translating cats to dogs or drawing trees into the sky). This is achieved by doing semantic segmentation on the source image, then translating that image and doing semantic segmentation on the target image. The loss describes the difference between the two resulting label maps. Formally they begin by learning a source model f_X that does semantic segmentation on the source data. For K-way classification with cross-entropy loss this results in following loss for the classification with L_X being the source labels:

$$\mathcal{L}_{\text{task}}(f_X, X, L_X) = -\mathbb{E}_{(x, l_X) \sim (X, L_X)} \sum_{k=1}^K \mathbb{1}_{[k=l_X]} \log(\sigma(f_X^{(k)}(x))) \quad (4.5)$$

With this the semantic loss is defined as:

$$\mathcal{L}_{\text{sem}}(G, F, X, Y, f_X) = \mathcal{L}_{\text{task}}(f_X, F(Y), p(f_X, Y)) \quad (4.6)$$

$$+ \mathcal{L}_{\text{task}}(f_X, G(X), p(f_X, X)) \quad (4.7)$$

with $p(f_X, x)$ being the label predicted by source semantic segmentation model f_X on a sample $x \in X$. Together with Equation 4.3 this results in the full objective of the CyCADA approach:

$$\mathcal{L}_{\text{CyCADA}}(f_Y, X, Y, L_X, G, F, D_X, D_Y) \quad (4.8)$$

$$= \mathcal{L}_{\text{task}}(f_Y, G(X), L_X) \quad (4.9)$$

$$+ \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \quad (4.10)$$

$$+ \mathcal{L}_{\text{cyc}}(G, F, X, Y) + \mathcal{L}_{\text{sem}}(G, F, X, Y, f_X) \quad (4.11)$$

And ultimately optimizing for a target model f_Y with:

$$f_Y^* = \arg \min_{f_Y} \min_{F, G} \max_{D_X, D_Y} \mathcal{L}_{\text{CyCADA}}(f_X, X, Y, L_X, G, F, D_X, D_Y) \quad (4.12)$$

SG-GAN [LLJX18] expands this by including a gradient-sensitive objective that emphasizes semantic boundaries by regularizing the generator to render distinct color/texture for each semantic region, and a patch-level discriminator that better understands differences in appearance distributions of different semantic regions (e.g. coarse texture of asphalt vs smooth and reflective of a vehicle). Prior to [ZPIE17] some works used paired images of the source and the target domain to train the networks. For example edgedrawings of shoes with a picture of the same shoe (see [IZZE16]). The cycle-consistency loss of CycleGAN makes it possible use datasets without paired images. This makes data acquisition easier and reduces costs which results in more training data and therefore makes better models possible. As CyCADA and SG-GAN are both based on CycleGAN they aswell don't require paired images.

5. Experiments

In this chapter the three Domain Adaptation Techniques *Cycle-consistent Generative Adversarial Network* [ZPIE17], *Cycle Consistent Adversarial Domain Adaptation* [HTP⁺17] and *Semantic-aware Grad-GAN for Virtual-to-Real Urban Scene Adaption* [LLJX18] will be compared by analysing similarities and differences. Furthermore pretrained models of each architecture were used to generate translated images on which a pretrained deeplabv3 model then was used to perform semantic segmentation and the results will be compared on their Intersection over Union scores.

5.1. Datasets

5.1.1. Synthetic dataset:

Playing for Data: Ground Truth from Computer Games

The GTA5 (Grand Theft Auto V) dataset is proposed in [RVRK16]. It contains 24966 images taken from a street view in the game Grand Theft Auto V by Rockstar Games [GTA]. The images are provided with 1914×1052 pixels and are containing moving cars, objects, pedestrians, bikes, have changing lighting and weather conditions aswell as day and night scenes. For all of these images the authors provide label images that are compatible with those of the Cityscapes dataset [COR⁺16]. Detouring, i.e. injecting a wrapper between the game and the graphics hardware to log function calls and reconstruct the 3D scene is used. This enables a faster labeling process as objects in a scene can be assigned an object ID through which the assigned label is propagated to other images containing this same object. Due to being more realistic than other existing synthetic street view datasets (e.g. SYNTHIA [RSM⁺16]) the GTAV dataset is very popular for training machine learning models related to autonomous driving and is therefore used in this work.

5.1.2. Real dataset:

The Cityscapes Dataset for Semantic Urban Scene Understanding

The Cityscapes dataset [COR⁺16] is a large scale dataset containing car dashcam view images from 50 european cities. It includes 30 classes relevant for autonomous driving. The images include scenes in spring, summer and fall seasons and under different weather conditions. There are 5000 images provided together with fine annotations and 20000 together with coarse annotations. Due to the large amount of

labeled data from a dashvram view and the inclusion of scenes with different weather and lighting conditions this dataset is often used to train deep neural networks that are related to autonomous driving. Due to the popularity and the GTAV dataset containing compatible labeling maps, this work uses Cityscapes as the real dataset for the experiments.

5.2. comparison Benchmark

5.2.1. Intersection over Union (IoU)

The Intersection over Union is a metric often used to compare semantic segmentation methods. All of the compared techniques were evaluated using IoU. It follows the following formula:

$$\frac{\text{predicted Pixels} \cap \text{ground truth Pixels}}{\text{predicted Pixels} \cup \text{ground truth Pixels}}$$

where predicted Pixels are the pixels predicted for a specific class by the semantic segmentation model and ground truth Pixels are the Pixels containing the ground truth for that image. Usually there are multiple different classes to predict and therefore it is common to calculate the mean IoU (mIoU) over all images that have predictions. To compare how well classes themselves are predicted by a model, one can also calculate the class IoU (cIoU).

5.2.2. Methodology

For the comparison each technique was used to translate a sample of 500 images from the GTA dataset to the Cityscapes domain. For CycleGAN and SG-GAN each, the authors provided pre-trained models. For CyCADA pretranslated images are provided in the project repository [CyCa]. Due to CyCADA only providing 22k translated images instead of the full 25k images provided in the GTA dataset, from the randomly chosen 500 images only 450 are available here. This has to be regarded in the following comparison. To translate images with SG-GAN and CycleGAN the provided code [SG],[Cycb] and for the semantic segmentation task an implementation [DLR] of deeplabv3 [CPSA17] was used. To compute the IoU values the deeplabv3 implementation uses the benchmark code provided by Cityscapes [CSR]. All computations were run on Ubuntu 18.04 using CPU only in VirtualBox on a Windows Machine due to problems with dependencies while running the code on the tcml cluster of the university of tuebingen [tcm], a cluster for machine learning applications.

6. Conclusion

To conclude...

A. Blub

Bibliography

- [COR⁺16] Marius Cordts, Mohamed Omran, Sebastian Ramos, Timo Rehfeld, Markus Enzweiler, Rodrigo Benenson, Uwe Franke, Stefan Roth, and Bernt Schiele. The cityscapes dataset for semantic urban scene understanding. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016.
- [CPSA17] Liang-Chieh Chen, George Papandreou, Florian Schroff, and Hartwig Adam. Rethinking atrous convolution for semantic image segmentation. *CoRR*, abs/1706.05587, 2017.
- [CSR] Cityscapes repository. <https://github.com/mcordts/cityscapesScripts>. Accessed: 2019-26-07.
- [Csu17] Gabriela Csurka. Domain adaptation for visual applications: A comprehensive survey. *CoRR*, abs/1702.05374, 2017.
- [CyCa] CyCADA repository. https://github.com/jhoffman/cycada_release. Accessed: 2019-01-08.
- [Cycb] CycleGAN lua repository. <https://github.com/junyanz/CycleGAN>. Accessed: 2019-01-08.
- [Cycc] Cycles rendering engine. <https://www.cycles-renderer.org/>. Accessed: 2019-10-07.
- [DLR] Deeplabv3 repository. <https://github.com/fregu856/deeplabv3>. Accessed: 2019-26-07.
- [Goo17] Ian J. Goodfellow. NIPS 2016 tutorial: Generative adversarial networks. *CoRR*, abs/1701.00160, 2017.
- [GPAM⁺14] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 2672–2680. Curran Associates, Inc., 2014.
- [GTA] Grand Theft Auto V website. <https://www.rockstargames.com/V/>. Accessed: 2019-26-07.

Bibliography

- [HTP⁺17] Judy Hoffman, Eric Tzeng, Taesung Park, Jun-Yan Zhu, Phillip Isola, Kate Saenko, Alexei A. Efros, and Trevor Darrell. Cycada: Cycle-consistent adversarial domain adaptation. *CoRR*, abs/1711.03213, 2017.
- [IZZE16] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. Image-to-image translation with conditional adversarial networks. *CoRR*, abs/1611.07004, 2016.
- [LLJX18] Peilun Li, Xiaodan Liang, Daoyuan Jia, and Eric P. Xing. Semantic-aware grad-gan for virtual-to-real urban scene adaption. *CoRR*, abs/1801.01726, 2018.
- [MPPS16] Luke Metz, Ben Poole, David Pfau, and Jascha Sohl-Dickstein. Unrolled generative adversarial networks. *CoRR*, abs/1611.02163, 2016.
- [Nas50] John F. Nash. Equilibrium points in n-person games. *Proceedings of the National Academy of Sciences*, 36(1):48–49, 1950.
- [RSM⁺16] German Ros, Laura Sellart, Joanna Materzynska, David Vazquez, and Antonio Lopez. The SYNTHIA Dataset: A large collection of synthetic images for semantic segmentation of urban scenes. In *CVPR*, 2016.
- [RVRK16] Stephan R. Richter, Vibhav Vineet, Stefan Roth, and Vladlen Koltun. Playing for data: Ground truth from computer games. In Bastian Leibe, Jiri Matas, Nicu Sebe, and Max Welling, editors, *European Conference on Computer Vision (ECCV)*, volume 9906 of *LNCS*, pages 102–118. Springer International Publishing, 2016.
- [SG] SG-GAN repository. <https://github.com/Peilun-Li/SG-GAN>. Accessed: 2019-01-08.
- [tcm] tcml cluster universität tübingen. <https://uni-tuebingen.de/fakultaeten/mathematisch-naturwissenschaftliche-fakultaet/fachbereiche/informatik/lehrstuehle/kognitive-systeme/projects/tcml-cluster/>. Accessed: 2019-01-08.
- [Uni] Unity graphics engine. <https://unity.com/>. Accessed: 2019-10-07.
- [Wu17] Jianxin Wu. Introduction to convolutional neural networks. *National Key Lab for Novel Software Technology. Nanjing University. China*, pages 5–23, 2017.
- [YXC⁺18] Fisher Yu, Wenqi Xian, Yingying Chen, Fangchen Liu, Mike Liao, Vashisht Madhavan, and Trevor Darrell. BDD100K: A diverse driving video database with scalable annotation tooling. *CoRR*, abs/1805.04687, 2018.
- [ZPIE17] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. *CoRR*, abs/1703.10593, 2017.